

# MDP

We are given a set of disconnected substations  $\mathcal{C}$ , with cardinality  $|\mathcal{C}| = N$  (in practice in Trieste they are always  $< 20$ ), between two remotely controlled substations, where these last will not be included in the problem, since they are already reconnected.

We define the cost of the process as the time each underlying user of each substation remains disconnected. So we compute the cost multiplying the time of disconnection for the number of users of a substation, and we sum them. We want to minimize this cost.

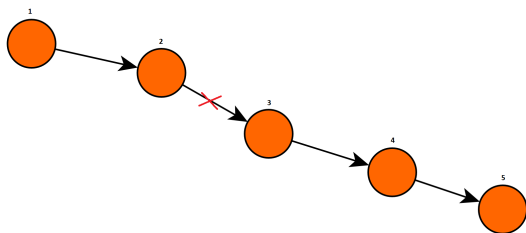


Figure 1. The fault has just occurred. All the substations are disconnected (orange).

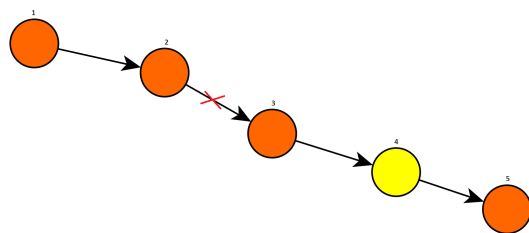


Figure 2. We visit substation 4 (yellow).

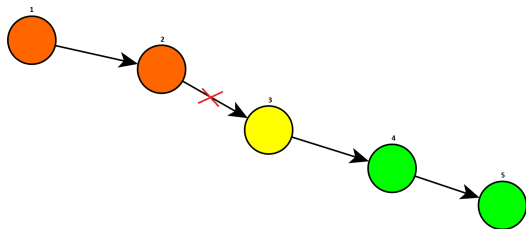


Figure 3. We have reconnected substations 4 and 5 (green). We go in substation 3 (yellow).

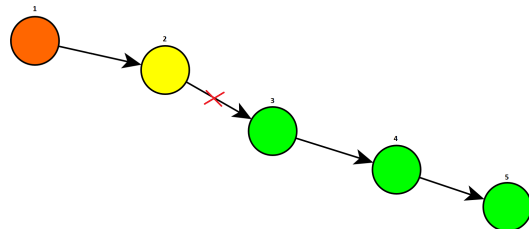


Figure 4. We have riconnected substation 3 (green). We are in substation 2 (yellow).




Figure 5. We reconnected substations 1 and 2 (green). All the subsations are reconnected.

In this MDP we have that

- the **state** is  $s = (x_g, v_k, \{v\})$  where  $x_g$  is the position of the fault,  $v_k \in \mathcal{C}$  is the substation in which I am, and  $\{v\}$  is the set of the still disconnected substations (we could also use the substation already reconnected, since they are complementary: given one of the two sets I can always retrieve the other). We have that  $v_k$  and  $\{v\}$  are *observable*, while  $x_g$  is *hidden*.  
**PROBLEM:** When the fault occurs the technician can be everywhere (at home if it is night, at the company, be around, etc.). The first state then what should it be? Do we put an extra fake substation that is the position of the technician? So at the beginning  $s_0 = (x_g, 0, \mathcal{C} = \{1, 2, 3, 4, 5\})$  and then for example  $s_1 = (x_g, 4, \{1, 2, 3\})$ . Is this good? The difference is that we use as  $\{v\}$  the set of disconnected substations **before** I operate on the substation in which I am or **after** I operate in the substation in which I am. If we use a substation "0" we consider as part of the problem optimizing going from it to another substation. If we don't have a "0" substation there are different initial states  $s_{01} = (x_g, 1, \mathcal{C} = \{1, 2, 3, 4, 5\})$ ,  $s_{02} = (x_g, 2, \mathcal{C} = \{1, 2, 3, 4, 5\})$ , ... But in this case what is the terminal state? [See Example]
- the **action** is the intervention I do in the specific substation I decide to visit, so  $a \in \mathcal{C}$ . Actually, since I can visit only disconnected substations, we have that  $a \in \{v\}$ .
- the **next state** is  $s' = (x_g, v_{k+1} = a, \{v'\})$  with  $\{v'\} \subseteq \{v\} \setminus a$ .
- the **reward** is the *cost* of going in a certain substation (as the time *in seconds* that it takes to go there from where I am) multiplied for the number of users still disconnected. Let's define as  $d_{v_k, v_{k+1}}$  the time *in seconds* to go from the substation  $v_k$  to the next substation  $v_{k+1}$ , and  $n_k$  the number of users still disconnected **before** operating in substation  $v_{k+1}$ . So if we are in state  $s = (x_g, v_k, \{v\})$ , we make an action and end up in state  $s' = (x_g, v_{k+1} = a, \{v'\})$ , we have that the number of disconnected users is  $n_k = \sum_{v \in \{v\}} u_v$ , where  $u_v$  is the number of users under substation  $v$ . So the reward is

$$r(s = (x_g, v_k, \{v\}), a, s' = (x_g, v_{k+1} = a, \{v'\})) = d_{v_k, v_{k+1}} \cdot n_k = d_{v_k, v_{k+1}} \cdot \sum_{v \in \{v\}} u_v \quad (1)$$

**Example 1:** Looking at the Figures 1-5, we have that the sequence of states and actions is:

- $s_0 = (x_g, 0, \mathcal{C} = \{1, 2, 3, 4, 5\}), a = 4$
- $s_1 = (x_g, 4, \{1, 2, 3\}), a = 3$
- $s_2 = (x_g, 3, \{1, 2\}), a = 2$
- $s_3 = (x_g, 2, \emptyset)$

If we don't use the fake substation 0 and  $\{v\}$  is the set of disconnected substations before, the sequence would be:

- $s_0 = (x_g, 4, \mathcal{C} = \{1, 2, 3, 4, 5\}), a = 4$  (in the state we have the action before we take it ?!?!?)
- $s_1 = (x_g, 3, \{1, 2, 3\}), a = 3$
- $s_2 = (x_g, 2, \{1, 2\})$
- $s_3 = (x_g, ?, \emptyset) \rightarrow s_3 = (x_g, 2, \emptyset)$  seems wrong

Since we know every aspect of the model of the environment, this is a **model-based** problem.

The system is **deterministic**, so given an admissible action  $a$  we will surely perform it and end up in the state in which that action leads. We can say that there are no execution errors, and I will always do what I want to do. This means that  $s' = \sigma(s, a)$ . So, mathematically we have that the **transition probability** is

$$p(s' | s, a) = \mathbb{I}(s' = \sigma(s, a)) = \delta_{s', s+a} = \begin{cases} 1 & \text{if } s' = \sigma(s, a) \\ 0 & \text{otherwise} \end{cases} \quad (2)$$

where  $\delta$  is the Kronecker delta and  $\mathbb{I}$  is the characteristic function of a set.

In our case, the transition probability is 1 only when, starting from state  $s = (x_g, v_k, \{v\})$ , the new substation  $v_{k+1}$  of  $s' = (x_g, v_{k+1}, \{v'\})$  is equal to the action  $a \in \mathcal{C} \setminus \{v\}$  that we took:

$$p(s' | s, a) = \begin{cases} 1 & \text{if } v_{k+1} = a \\ 0 & \text{if } v_{k+1} \neq a \end{cases} \quad (3)$$

This is a problem of partial **conservabilità** since part of the state  $s$  is hidden (the position of the fault  $x_g$ ). So, we can not solve it using dynamic programming, because in this way the policy would depend on the whole state, thus also on the hidden state, which can not be.

If we know where the fault is, the solution is straightforward: if it is on an edge we visit the two substations at the ends of it (we might have to choose the right order), if it is in a substation we visit it directly.

If we don't know the position of the fault  $x_g$ , we can solve this problem with two different approaches. One is to use **approximate value iteration**: we work in a subspace of the states space corresponding to the observable states. Another possibility is to perform **policy iteration**, or gradient descent in the policy space. So we impose that the policy depends only on the observable variables, and we try to find the best policy in this subspace in order to optimize the MDP.

So since we don't know where the failure is, the **policy** depends only on the observable states, so it doesn't know where the fault is. Let's define a parameterized policy using Boltzmann parameterization:

$$\pi(a | s = (x_g, y = (v_k, \{v\}))) = \frac{e^{\theta_a y}}{\sum_b e^{\theta_b y}} \quad (4)$$

where  $\theta$  are the parameters which depends only on the observable variable  $y$ .

If we search in this space of policies, this will give us a policy which doesn't depend on the time, since with any algorithm we try to find the optimal parameters to solve this problem. This gives us a **stationary policy**. The structure of the states is already a measure of time, since we have the number of steps already done: the substations we already visited. So the important is not to establish a policy at the different steps, but a policy with respect to the states.

This is a problem with terminal state, which occurs when I reconnect all the substations. It will always be reached, since with every action I visit a substation, and at the very least I remove it from the set of disconnected substations (if I'm lucky I can remove half of the substations from the set of disconnected substations every time). So, for this specific problem it doesn't make sense to introduce a discount factor  $\gamma$ .

Let's define  $J$  as the sum of all the costs you have if you are in state  $s = (x_g, v_k, \{v\})$  summed in time until the process is concluded. The steps of the process are formal steps, since in a step we pass from one substation to another, so the physical time is in the costs (as the time / cost of going from one substation to another).

So we associate a *cost function*  $J_\pi$  to each state  $s \in \mathcal{S}$  that is the *accumulated cost from that state to the end of the process, following policy  $\pi$* . So we have that the cost is

$$\begin{aligned} J_\pi(s) &= \mathbb{E}_\pi \left[ \sum_{t=0}^{\infty} r(s_t, a_t, s_{t+1}) \right] \\ &= \mathbb{E}_\pi \left[ \sum_{k=0}^{\infty} d_{v_k, v_{k+1}} \cdot n_{k+1} \mid s_0 = s \right] \end{aligned} \quad (5)$$

If we can find an optimal path in the states, we can find an optimal path in the substations, since that state contains the current substation and other information.

*Can we find a recursive relation like in the Traveling Salesman?*

$$J_\pi(s = (x_g, v_k, \{v\})) = \sum_{a \in \{v\}} \pi(a|s) \left[ d_{v_k, v_{k+1}} \cdot n_{k+1} + J_\pi(s' = (x_g, a, \{v'\})) \right] \quad (6)$$

(from the recursive equation of  $V$ , since  $V = J$ ). Since we don't know which will be the last visited substation, the terminal state is not unique: we have  $\mathcal{C}$  terminal states, which are  $s = (x_g, v_k, \emptyset)$  with  $v_k \in \mathcal{C}$ . Actually, we have that  $v_k$  must be *one of the two substations near the fault, or the substation in which we have the fault*. So there are at most two terminal states, possibly even only one. Since the transition probabilities are deterministic, from each one of them we can go back. But unlike in the Traveling Salesman, we won't solve the Bellman's equation at each step performing a  $\max$  on the actions, since the policy can not depends on the whole state, but only on the observable part.

This is why we have to perform a continuous projection in the case of approximate value iteration or the gradient search. The approximate value iteration is a delicate operation and you have to pay a lot of attention, instead using the gradient is more direct. So we will use the latter.

## Gradient descent in the policy space

We have that the gradient formula is

$$\nabla_\theta J = \sum_{s,a} \eta_\pi(s) Q_\pi(s, a) \nabla_\theta \pi(a|s). \quad (7)$$

To optimize  $J$  we perform a gradient descend on  $\theta$ . The quantities  $\eta$  and  $Q$ , given a policy  $\pi$ , are computed through *linear* operations (since they obey to linear equations), and are computed for all the states. They don't contain the position of the fault, then we sum over all possible positions of the fault.

**Initial state:** I didn't visit any substation and I have all the possible positions of the fault.

**Idea for the state:** We could use as state everything that is included in the ordered pair of two substations:  $\{v\} = (v_l, v_r)$ . And every intervention sends me from one state to another.

Topological assumption: we have a tree structure. We have to find a way of representing forks.

The theory says that

$$\begin{aligned} Q_\pi(s, a) &= \sum_{s'} p(s'|s, a) \left( r(s, a, s') + \sum_{a'} \pi(a'|s') Q(s', a') \right) \\ &= \mathbb{E} \left[ \sum_{t=0}^{\infty} r(s_t, a_t, s_{t+1}) \mid s_0 = s, a_0 = a \right] \end{aligned} \quad (8)$$

is the **state-action value function** or the **quality** of the state-action pair. Given that the state is  $s = (x_g, v_k, \{v\})$ , the action is  $a \in \{v\}$  and the new state is  $s' = \sigma(s, a) = (x_g, v_{k+1} = a, \{v'\})$ , the equation of  $Q$  in our case, given (3), is:

$$Q_\pi(s, a) = \left( d_{v_k, a} \cdot n_{k+1} + \sum_{a' \in \{v'\}} \pi(a' | \sigma(s, a)) Q((\sigma(s, a), a')) \right) \quad (9)$$

While for the other variable we have that

$$\eta_\pi(s') := \rho_0(s') + \sum_{s, a} \pi(a | s) p(s' | s, a) \eta_\pi(s) \quad (10)$$

is **the time spent in state  $s'$  before the process dies**, where  $\rho_0(s')$  is the probability of starting in the initial state  $s'$ .

We can notice that the value of  $Q$  depends on the values of  $Q$  of future states, while the value of  $\eta$  depends on the values of  $\eta$  for previous states.

In (8) we have that for every value of the parameters  $\theta$  we know  $\pi(a' | s')$ , so it is a linear equation in  $Q$ .

We have that that states are in the order of the number of substations cubed  $N^3$ , since we have the position of the fault and the two substations that identify the not yet visited substations, and the actions are in the order of  $N$ . So  $Q$  is an equation in  $N^4$  variables. We solve it iteratively using value iteration (we find the fixed point).

I can go from any substation to any other substation that has not already been connected. Let's suppose that my policy is simply to go in a random substation still disconnected, so  $\pi$  is a uniform distribution over the disconnected substations, so it is

$$\pi(a | s = (x_g, v_k, \{v\})) = \frac{1}{|\{v\}|} \quad (11)$$

So since  $s' = (x_g, v_{k+1} = a, \{v'\})$  we have that  $Q$  becomes

$$Q_\pi(s = (x_g, v_k, \{v\}), a) = d_{v_k, a} \cdot n_{k+1} + \sum_{a' \in \{v'\}} \frac{1}{|\{v'\}|} Q(\sigma(s, a), a') \quad (12)$$

Let's try to estimate how much does it cost computing  $Q$ . If we use the QR factorization for rectangular matrices on  $Q \in \mathbb{R}^{|\mathcal{S}| \times |\mathcal{A}|}$ , it costs  $O(|\mathcal{S}| \times |\mathcal{A}|^2)$ . Since we have that the number of substations is  $|\mathcal{C}| = N$ , given that a state can be thought as the current substation  $v_k$  and the pair of substations that delimit the segment of disconnected substation, we have that  $|\mathcal{S}| \sim O(N \cdot N^2) = O(N^3)$ , while  $|\mathcal{A}| \sim O(N)$ . So we have that solving this linear system requires  $O(|\mathcal{S}| \times |\mathcal{A}|^2) = O(N^3 \times N^2) = O(N^5)$ . **CHECK WHAT PROF SAID ABOUT SPEED**

We have that in (14) the function  $\rho_0(s')$  is the probability of starting in the initial state  $s'$ . Since we don't have any prior information of where the fault might be,  $\rho_0$  doesn't depend on it, so  $\rho_0$  will be uniform in  $x_g$  (this means that we divide by the number of possible positions of  $x_g$ , which is either a substation or an edge, so it is  $N + (N + 1) = 2N + 1$ ). Since the first substation in which to go is random, it doesn't depend on  $v_k$  either. Instead, the set of the disconnected substations must be equal to the set of all the substations  $\mathcal{C}$ . So  $\rho_0$  must be 1 when the set of the disconnected substations is equal to  $\mathcal{C}$  and must be 0 for every other proper subset. So we have that

$$\rho_0\left(s = (x_g, v_k, \{v\})\right) = \frac{\star}{|\mathcal{C}|(2|\mathcal{C}| + 1)} \mathbb{I}(\{v\} = \mathcal{C}) = \frac{\star}{N(2N + 1)} \mathbb{I}(\{v\} = \mathcal{C}) \quad (13)$$

since the number of possible  $x_g$  is  $|\mathcal{C}|$  (not true, the fault can be also on the edges, not only in the substations!) and the number of possible initial substations is  $|\mathcal{C}|$  (what about the fact that we start from where the technician is at the moment of the fault, and then he goes in the first substation?).

If we have prior information on substations that have frequent faults, we can change this probability distribution to use this information.

So in our case, thanks to (13), (3), and (11), we have that

$$\begin{aligned} \eta_\pi\left(s' = (x_g, v_{k+1}, \{v'\})\right) &:= \rho_0(s') + \sum_{s,a} \pi(a|s) p(s'|s, a) \eta_\pi(s) \\ &= \frac{1}{N^2} \mathbb{I}(\{v'\} = \mathcal{C}) + \sum_s \pi(a = v_{k+1}|s) \eta_\pi(s) \\ &= \frac{1}{N^2} \mathbb{I}(\{v'\} = \mathcal{C}) + \sum_s \frac{1}{|\{v\}|} \eta_\pi\left(s = (x_g, v_k, \{v\})\right) \end{aligned} \quad (14)$$

**Idea of the algorithm:** We start from a certain policy, for example the random policy of equation (11), in which we choose randomly the substation to be visited. This means that in the parameterized policy (4) all the parameters  $\theta$  are equal to 0. This is because in the random policy the parameters  $\theta$  don't depend on the action  $a$ , so we have that

$$\pi\left(a \mid s = (x_g, y = (v_k, \{v\}))\right) = \frac{e^{\theta y}}{\sum_{b \in |A|} e^{\theta y}} = \frac{e^{\theta y}}{e^{\theta y} \sum_{b \in \{v\}} 1} = \frac{1}{|\{v\}|}. \quad (15)$$

This is true for every choice of  $\theta$  which doesn't depend on the action, but in practice to construct a uniform policy we take  $\theta = 0$ . But notice that this is not the only way.

In general, there is no guarantee that this is a convex problem in the parameters  $\theta$ , we could have several different minima. So one should do random restarts with different policies than the random one to see if you can reach a different minimum. However, this doesn't guarantee to find the global minimum, but it is the best we can do.

Then we compute the gradient using (7). The gradients of the policy are simple computations (since we decided the parametrization of the policy, we can derive it), while the objects  $Q$  and  $\eta$  have to be done indirectly: you have to solve the linear equations (12) and (14) for that given policy. This can be more or less computationally heavy, but we have to solving two linear equations, with any chosen method. Given the form of our transition probabilities  $p(s'|s, a)$ , we have that the equations of  $Q$  don't depend on  $s'$ , but only on  $a$ . So fixed the first state  $s$ , we have to solve two linear equations only on the actions (we can see the  $Q(s, a)$  matrix as a series of vectors). **This operations can be parallelized!**

So having done these steps we have an expression for the gradient for every parameter. So we do a sweep over all the parameters and we find the value of the gradient for each one of them. Then we take a step in the parameters space and we descent the gradient. We stop when the value of  $J$  doesn't change "so much" in percentage.

If we have the intuition that there is a deterministic sequence of action to be performed, we have that the parameters  $\theta$  tend to infinity. This is because the deterministic policies correspond to a one-hot vector in which we have 1 for only one action and 0 for all the other actions, and this happens when the parameter associated to this action becomes much bigger than the others, so that when we normalize it with the parametrization of the policy (which is like a softmax) it becomes 1 while the others are so small that becomes 0. If it is like this, the gradient descent never stops and there could be problems of overflow on  $\theta$  since it goes to infinity.

A smart thing to do when this happens is to write the parametrization in the following way (the state  $s$  is fixed):

$$\pi(a) = \frac{e^{\theta_a}}{\sum_b e^{\theta_b}} = \frac{e^{-(\max_{a'} \theta_{a'} - \theta_a)}}{\sum_b e^{-(\max_{a'} \theta_{a'} - \theta_b)}} . \quad (16)$$

The benefit of writing it in this way is that, since the exponents of the two exponentials are negative (the parts in the parenthesis  $(\max_{a'} \theta_{a'} - \theta_a)$ ) are always positive), this never explodes. This is a simple numerical trick that allows to improve the stability of the algorithm.

## Model free

If this is not possible, one option is to do it model free, so instead of doing a complete gradient we do a stochastic gradient, using the **policy gradient algorithms**. In this case we remove the cost of solving the equations for  $Q$  and  $\eta$ , since we proceed by trial and error, you try an action, compute a policy, upgrade the gradient and go on, and this operations are not computationally heavy: you just need to experience the costs. The downside is that it is stochastic, so it has a great variance.

## Natural policy gradient

Another option that can be done both at the level of the deterministic gradient and at the level of the stochastic gradients is to use **natural policy gradient**. It's a minimal change: we only change the gradient of  $\pi$ , everything else remains the same.

The improvement is that instead of having a gradient that reaches a horizontal asymptote, the gradient keeps on moving and doesn't slow down. This allows to reach the convergence faster.